

Searching the State Space


Part 1

State space

- A **state** is a data structure (object) that represents a possible configuration of the world (agent and environment).
- The **state space** of a problem is the set of all possible states for that problem.
- Example: A vacuum cleaner agent in two adjacent rooms which can be either clean or dirty.

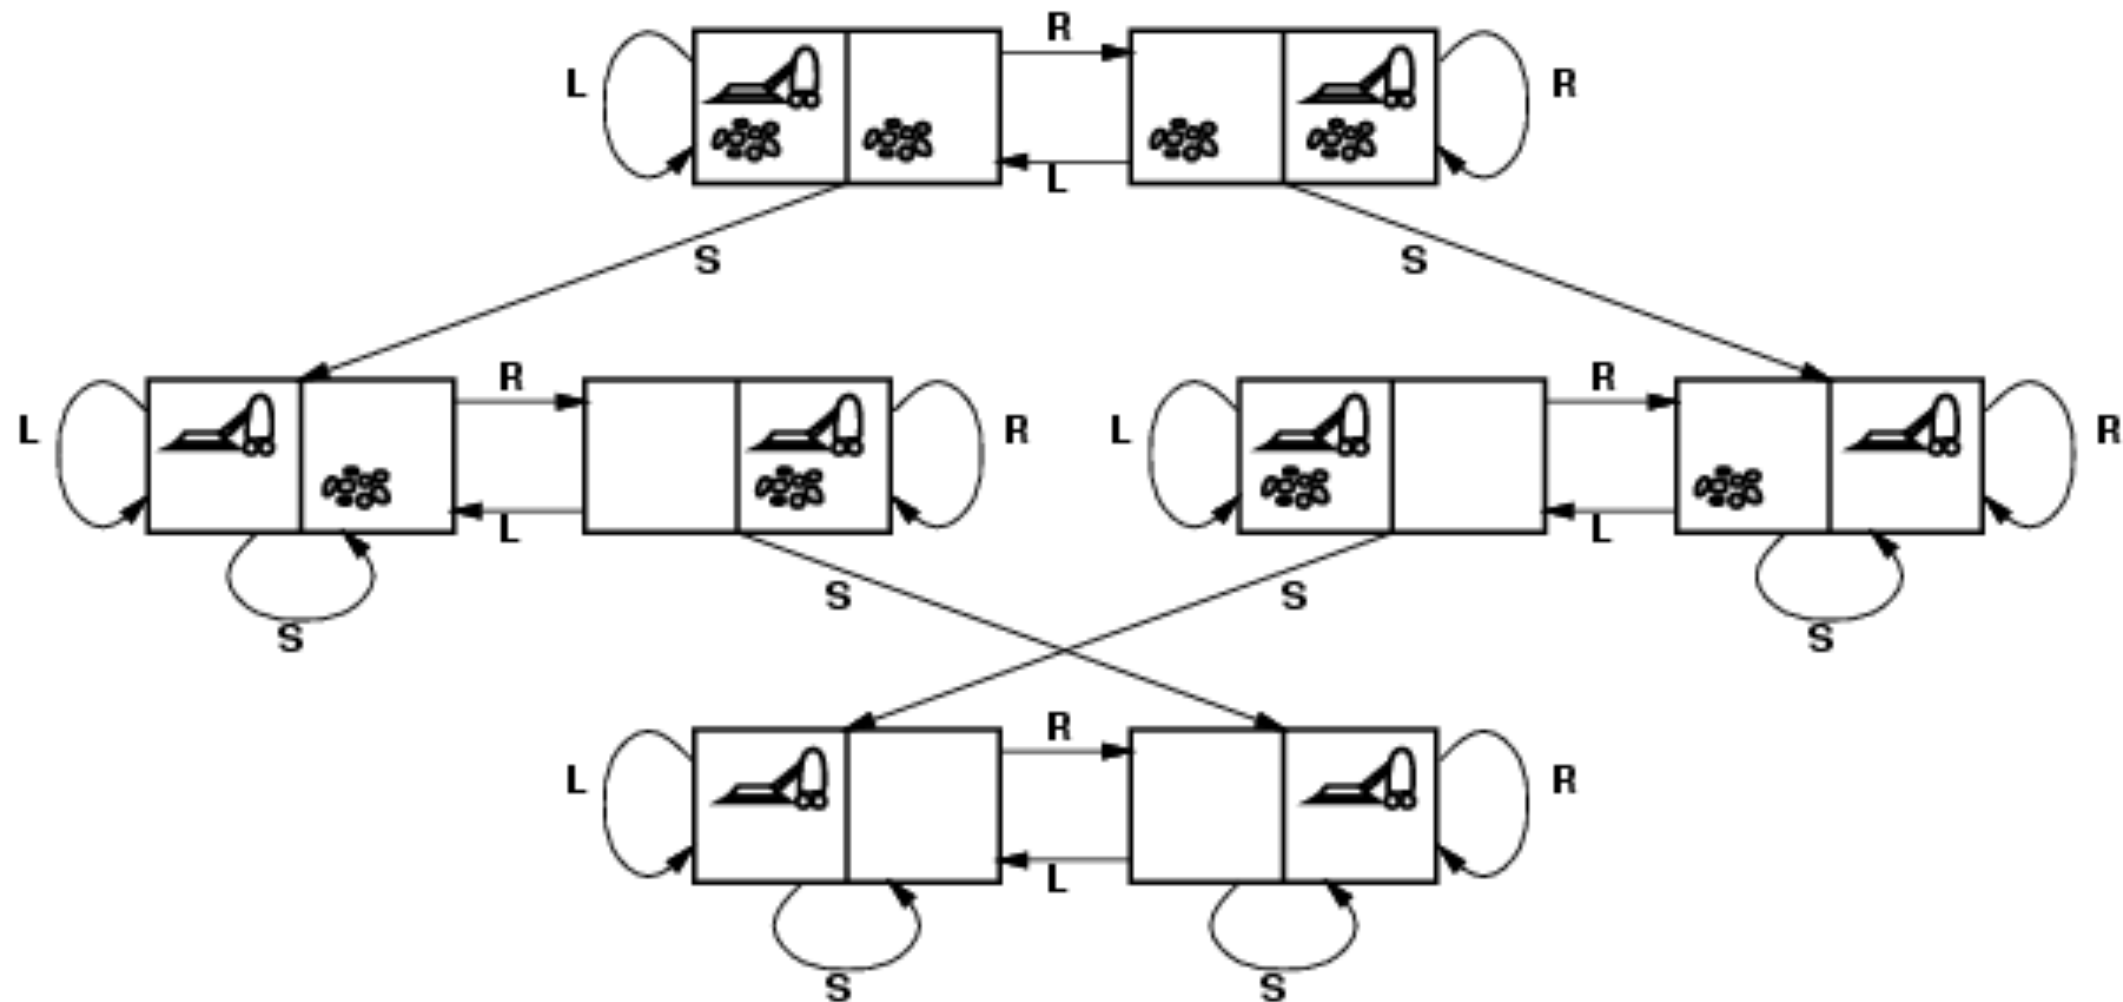
- Location = {left, right}
- Left-room-condition = {dirty, clean}
- Right-room-condition = {dirty, clean}
- State-space = Location
× Left-room-condition
× Right-room-condition

In this example, each state is represented by a triple (3-tuple).



State space graphs

- Actions change the state of the world.
- Example: Suppose the vacuum cleaner agent can take the following actions: L (go left), R (go right), S (suck).



Directed graphs

- Many AI problems can be abstracted into the problem of finding a path in a directed graph.
- A graph consists of a set of **nodes** and a set of **arcs**.
 - nodes are usually used for states.
 - arcs are used for actions.
- In theory, a **path** is sequence of nodes $\langle n_0, n_1, \dots, n_k \rangle$ where there is an arc from n_i to n_{i+1} . In practice, a path is sequence of arcs (see the provided Python module in the quiz.)
- The **length** of path $\langle n_0, n_1, \dots, n_k \rangle$ is k .
- An arc can have an associated cost (which could mean distance, time, energy, etc). The **cost** of a path is the sum of the cost of the arcs along the path.
- Given a set of **starting nodes** and a set of **goal nodes** (desired states), a **solution** is a path from a start node to a goal node.

Explicit vs Implicit graphs

- In **explicit graphs** nodes (vertices) and arcs (edges) are readily available. They are read from the input and stored in a data structure such as an adjacency list or an adjacency matrix.
 - The entire graph is in the memory
 - The complexity of algorithms are measured in terms of the number of vertices and/or edges.
- In **implicit graphs** a procedure **outgoing_arcs** (or *neighbours*) is defined that given a node, returns the set of directed arcs that connect the node to other nodes.
 - The graph is generated on the fly.
 - The complexity of algorithms are measured in terms of the *depth* of a goal node and the *average branching factor* (average number of outgoing arcs).

Example: the 8-puzzle

- States: location of tiles
- Actions: Move tiles that can be moved (or move the blank space) left, right, up, down.

7	2	4
5		6
8	3	1


Start State

	1	2
3	4	5
6	7	8

Goal State


Example: Pac-Man

- States: permutation of food and Pacman in its world.
- Actions: N, E, S, W (when possible)
- Goal states: eaten all the food



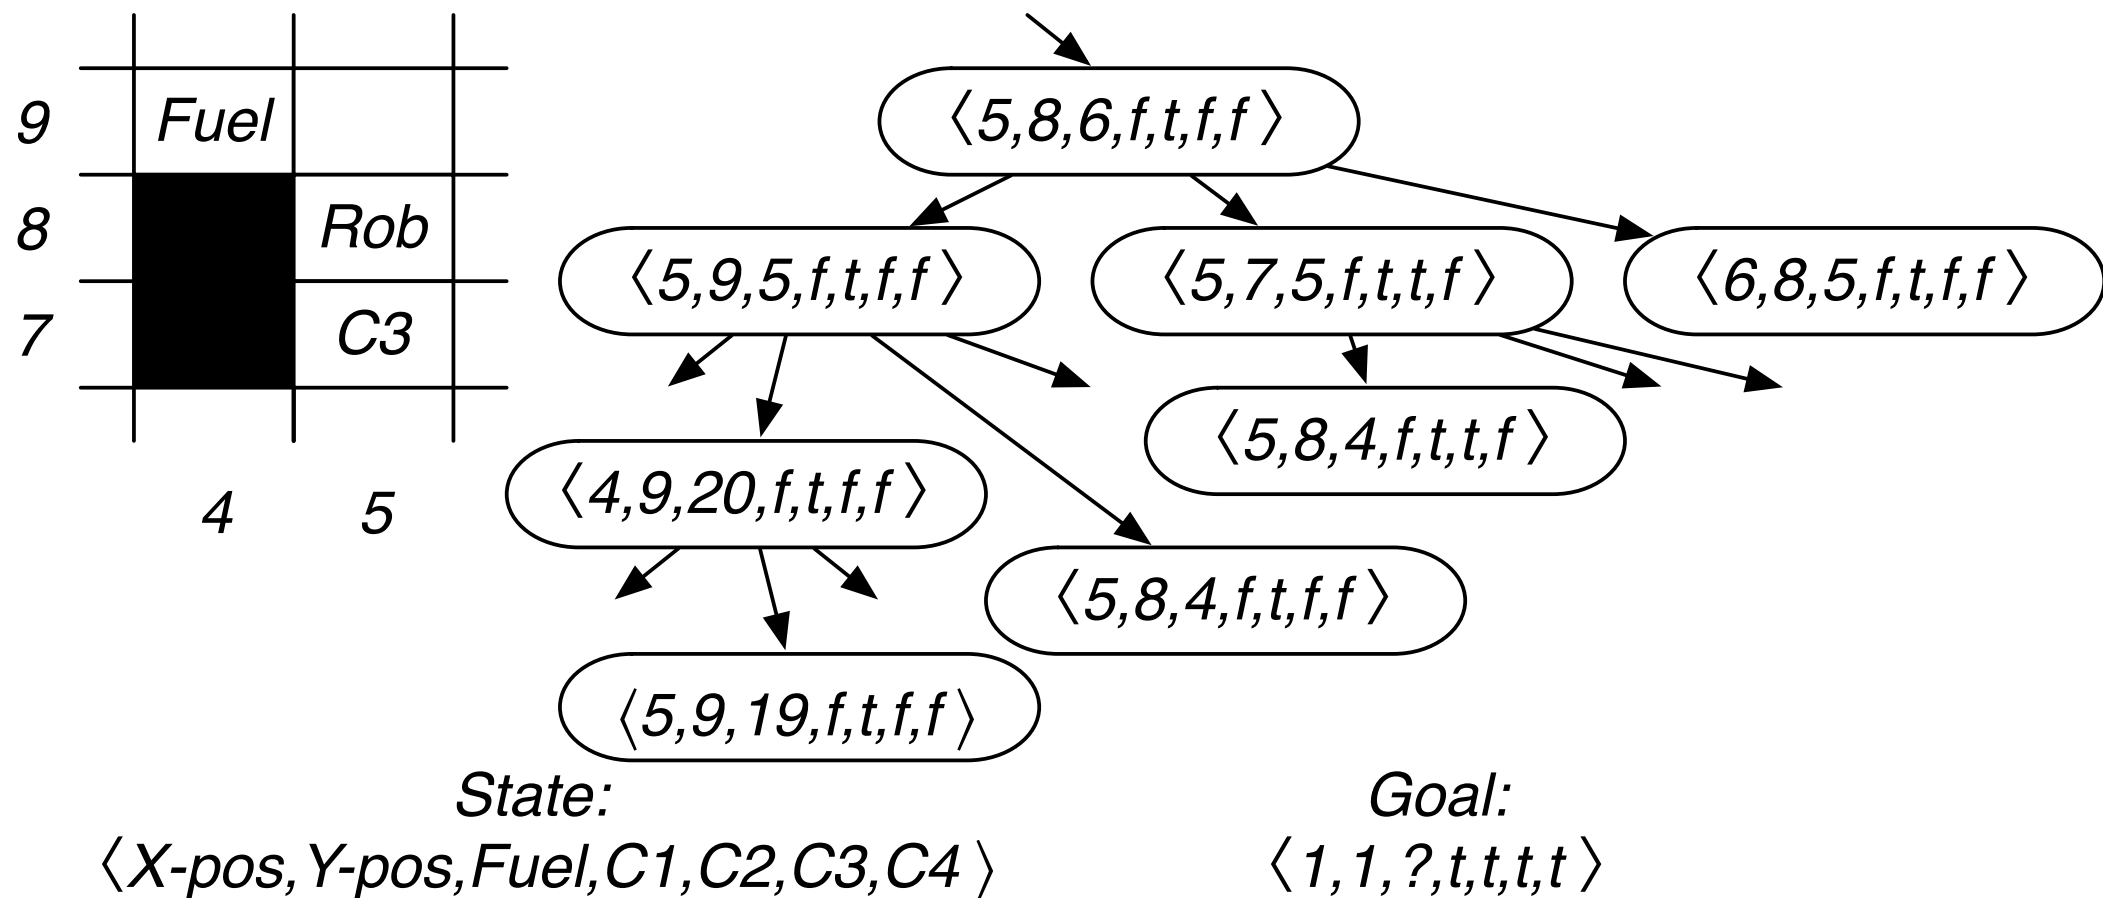
Example: holiday in Romania

- We are on holiday in Romania; currently in Arad. Flight leaves tomorrow from Bucharest. We need to get there.



Example: a grid game


- Grid game: Rob needs to collect coins C1, C2, C3, C4, without running out of fuel, and end up at location (1, 1)



Searching graphs

- Generic search algorithm: given a graph, start nodes, and goal nodes, incrementally explore paths from the start nodes.
- Maintain a **frontier** of paths that have been explored. [All the paths are from the starting node(s).]
- As search proceeds, the frontier is updated and the graph is explored until a goal node is encountered.
- The way (order) in which paths are removed from (and added to) the frontier defines the **search strategy**.

Search tree



Generic graph search algorithm

Input: a graph,

a set of start nodes,

Boolean procedure $goal(n)$ that tests if n is a goal node.

$frontier := \{ \langle s \rangle : s \text{ is a start node} \};$

while $frontier$ is not empty:

select and **remove** path $\langle n_0, \dots, n_k \rangle$ from $frontier$;

if $goal(n_k)$

return $\langle n_0, \dots, n_k \rangle$;

for every neighbor n of n_k

add $\langle n_0, \dots, n_k, n \rangle$ to $frontier$;

end while


Remarks

- Which value is selected from the frontier at each stage defines the search strategy. In our implementation we pass a frontier object to the search procedure. The frontier object is in charge of selecting the next path.
- The function *neighbors* defines the graph. In our implementation we use the method `outgoing_arcs` for this purpose.
- The function *goal* defines what is a solution. In our implementation we use the method `is_goal` for this purpose.
- If more than one answer is required, the search can continue. For this reason, in our implementation we use `yield` instead of `return`.

Depth-first search

- In order to perform DFS, the generic graph search must be used with a stack frontier (last in, first out).
- If the stack is a Python list of the form $[..., p, q]$ where each element is a path, then
 - q is selected (popped).
 - if the algorithm continues, then paths that extend q are pushed (appended) to the stack.
 - p is only selected (popped) when all paths from q have been explored.
- As a result, at each stage, the algorithm expands the deepest path.

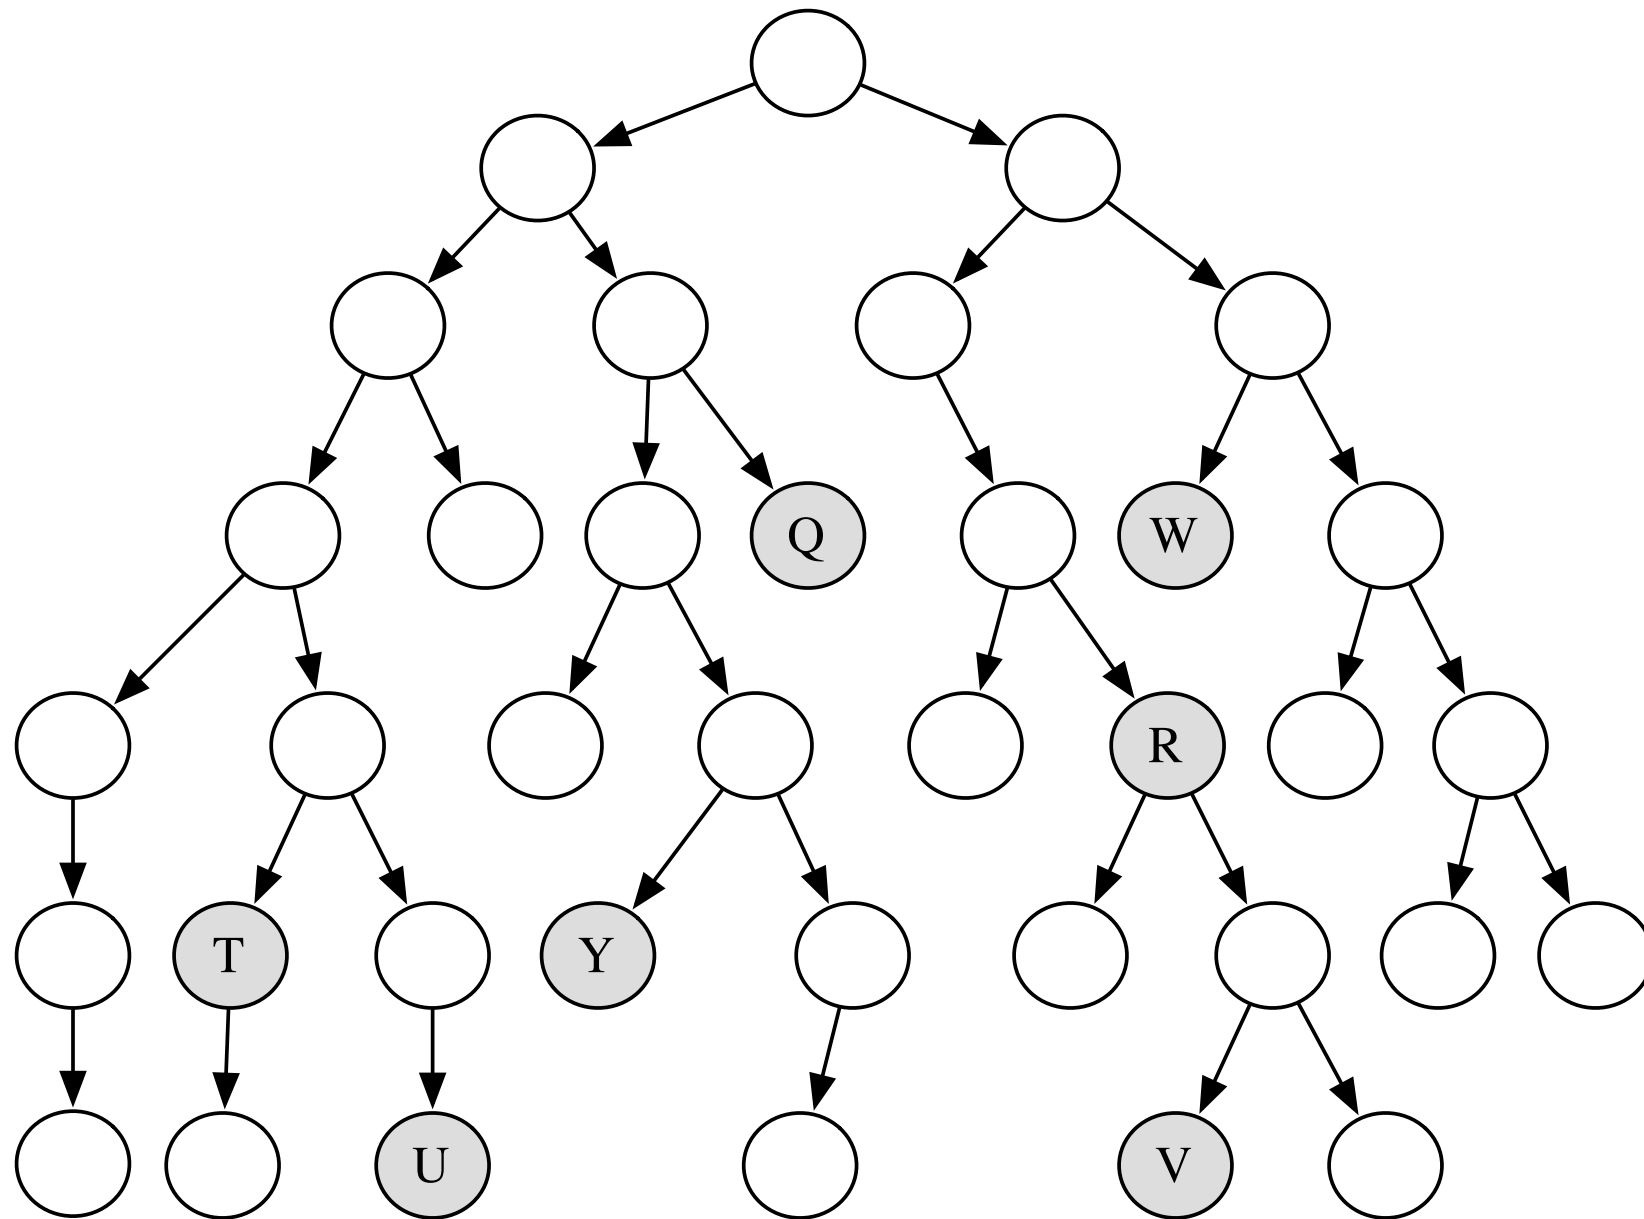
DFS: illustration



- [Assuming that in the sequence of Arcs returned by the function *neighbors* (or the method `outgoing arcs` in the implementation), the right arc comes before the left arc.]

DFS: behaviour

- Which shaded goal will a depth-first search find first?



DFS: behaviour and complexity

- Does DFS guarantee to find a solution with the fewest arcs (if there is a solution)?
- Is it complete? (i.e. does it guarantee to find a solution if there is one?)
- Does it halt on every graph?
- Does the goal state affect the search behaviour?

Assume a finite search tree of depth ***d*** and branching factor ***b***:

- What is the time complexity?
- What is the space complexity?

Explicit graphs in quizzes

- In some exercises we use small explicit graphs to study the behaviour of various search frontiers.
- Nodes are specified in a set. The use of set should remind you that the order of the nodes does not matter. The elements of the sets are strings (node names).

```
nodes={'a', 'b', 'c', 'd'}
```

- Edges are specified in a list. The use of list should remind you that the order of the edges matters. The elements are either:
 - ♦ pairs of nodes (tail, head); or
 - ♦ triples of nodes (tail, head, cost).

```
edge_list = [('a', 'b', 5), ('a', 'c', 3)]
```

How do we trace the frontier

- Starting with an empty frontier we record all the calls to the frontier: to add or to get a path. We dedicate one line per call.
- When we ask the frontier to add a path, we start the line with a **+** followed by the path that is being added.
- When we ask for a path from the frontier, we start the line with a **-** followed by the path that is being removed.
- When using a priority queue, the path is followed by a comma and then the key (e.g. cost, heuristic, f-value, ...).
- The lines of the trace should match the following regular expression (case and space insensitive): `^[+-][a-z]+(,\d+)?!?$`

Example: tracing frontier in DFS

Given the following graph

```
nodes={a, b, c, d},  
edge_list=[(a,b), (a,d), (a, c), (c, d)],  
starting_nodes = [a],  
goal_nodes = {d}
```

trace the frontier in depth-first search (DFS).

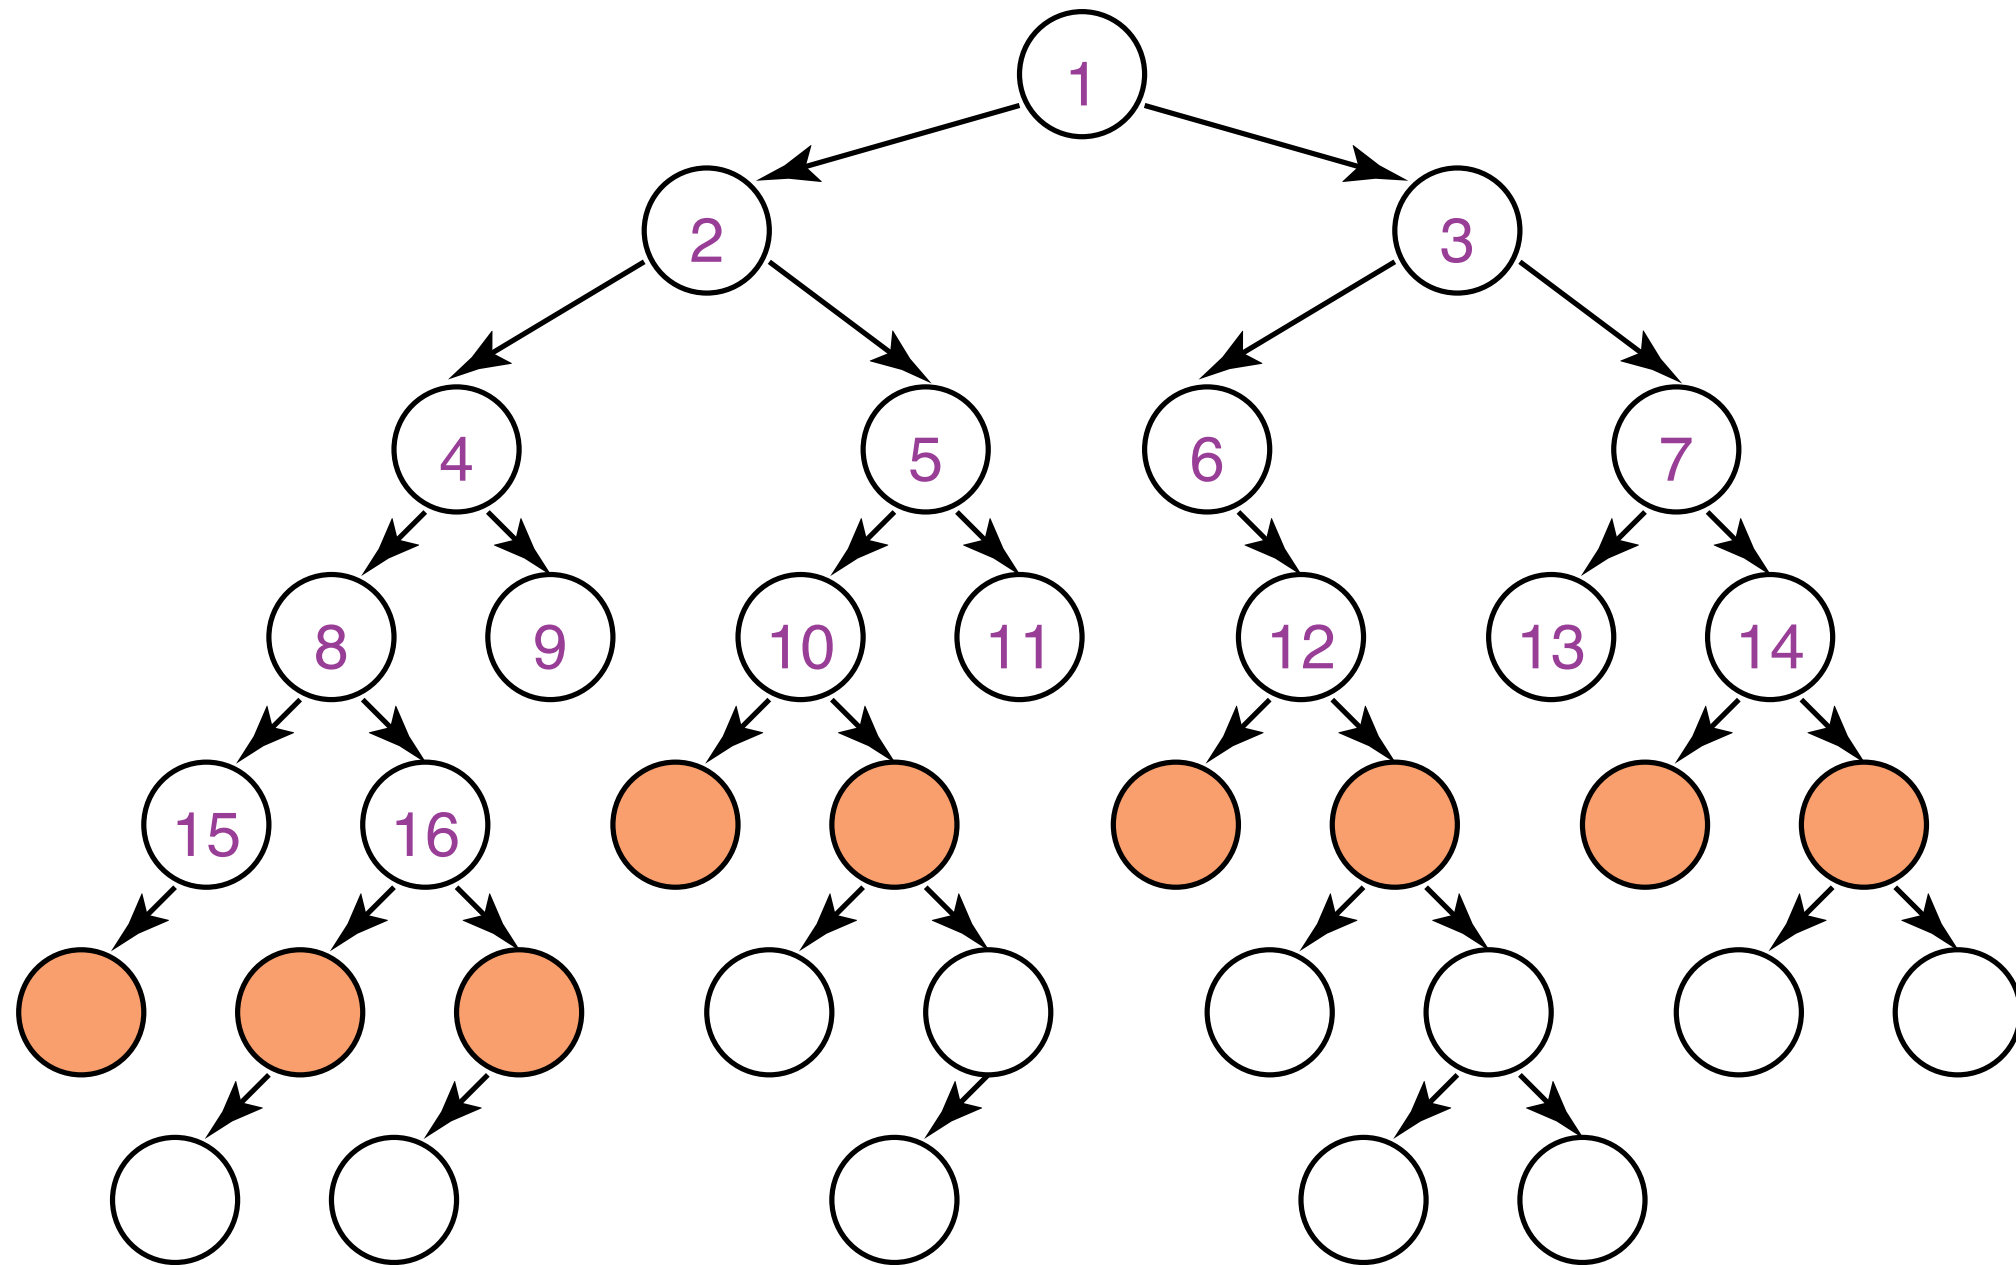
Answer:

```
+ a  
- a  
+ ab  
+ ad  
+ ac  
- ac  
+ acd  
- acd
```

Breadth-first search

- In order to perform BFS, the generic graph search must be used with a queue frontier (first in, first out).
- If the queue is a Python deque of the form $[p, q, \dots, r]$, then
 - p is selected (dequeued from left).
 - if the algorithm continues then paths that extend p are enqueued (appended) to the queue after r .
 - in the next iteration q is selected (dequeued) from the frontier.
- As a result, at each stage, the algorithm expands the shallowest path.


BFS: illustration of search tree



- [Assuming that in the sequence of Arcs returned by the function *neighbors* (or the method `outgoing_arcs` in the implementation), the left arc comes before the right arc.]

BFS: behaviour

- Which shaded goal will a breadth-first search find first?



BFS: behaviour and complexity

- Does BFS guarantee to find a solution with the fewest arcs (if there is a solution)?
- Is it complete? (i.e. does it guarantee to find a solution if there is one?)
- Does it halt on every graph?
- Does the goal state affect the search behaviour?

Assume a finite search tree of depth ***d*** and branching factor ***b***:

- What is the time complexity?
- What is the space complexity?

Example: tracing frontier in BFS

Given the following graph

```
nodes={a, b, c, d},  
edge_list=[(a,b), (a,d), (a, c), (c, d)],  
starting_nodes = [a],  
goal_nodes = {d}
```

trace the frontier in breadth-first search (BFS).

Answer:

```
+ a  
- a  
+ ab  
+ ad  
+ ac  
- ab  
- ad
```

Lowest-cost-first search (LCFS)

- The cost of a path is the sum of the costs of its arcs.
- LCFS selects a path on the frontier with the lowest cost.
- The frontier is a priority queue ordered by path cost.
- LCFS finds an optimal solution: a least-cost path to a goal node.
- Another name for this algorithm is *uniform-cost search* (which is somewhat misleading).

Priority queue refresher

1. A container in which each element has a priority (cost).
2. An element (path) with higher priority (lower cost) is always selected/removed/dequeued before an element with lower priority (higher cost).
3. We require the priority queue to be stable: if two or more elements have the same priority, elements that were enqueued earlier are dequeued earlier.
4. In Python you can use `heapq`. You need to store objects in a way that the above properties hold.

Priority queue example

On an empty frontier, after executing:

- + c, 5
- + b, 10
- + a, 5
- + d, 10

a sequence of selection/removals yields:

- c, 5
- a, 5
- b, 10
- d, 10

Example: tracing frontier in LCFS

Given the following graph

nodes={a, b, c, d, g},
edge_lists=[(a,b,4), (a,c,2), (a,d,1),
 (b,g,4), (c,g,2), (d,g,4)],
starting_nodes = [a],
goal_nodes = {g}

trace the frontier in lowest-cost-first
search (LCFS).

Answer:

+ a, 0
- a, 0
+ ab, 4
+ ac, 2
+ ad, 1
- ad, 1
+ adg, 5
- ac, 2
+ acg, 4
- ab, 4
+ abg, 8
- acg, 4